

# MET CS 767 Assignment 2: Neural Nets Intro

## *Brendan Torok*

### **1. How I modified data and/or code to attempt improvement**

#### 1.1 Description of your modifications and reason(s) it could be an improvement

For this assignment I used the base neural network provided in the example code and applied it to my custom dataset as a starting point for training my neural network for my chosen application.

My main changes are as follows:

Increased layers to 4 layers – this increased performance of the neural network although it seems to have made the neural network learn very quickly, within 1-2 epochs. This is an architectural change.

Added manual learning rate tuning to Adam optimizer. This greatly improved the performance of the model. This is a parametric change.

```
optimizer = tf.keras.optimizers.Adam(learning_rate=0.0004)
```

Changed activation functions to try gelu, and leaky activation. These both decreased performance so I kept the relu activation function. This is an architectural change.

Added L2 regularization and tuned regularization parameter to improve performance. This is both an architectural and parametric change.

```
tf.keras.layers.Dense(128, activation='relu',
                      kernel_regularizer=tf.keras.regularizers.l2(8e-4))
```

Adjusted dropout to 0.10 and 0.05 after each layer, changed from 0.2. This is a parametric change.

```
tf.keras.layers.Dropout(0.10),
tf.keras.layers.Dense(128, activation='relu',
                      kernel_regularizer=tf.keras.regularizers.l2(8e-4)),
tf.keras.layers.Dropout(0.05)
```

Added three more layers to the neural network, two with more neurons and one with the same number of neurons in the layer. This was an improvement and is an architectural change.

In training I adjusted the number of epochs and tested early stopping. I discovered that adding more epochs did not greatly improve my performance, but adding early stopping

caused the neural network training to stop after 1 or 2 epochs, resulting in poor performance. I settled with 20 epochs as it seemed to have the most stability.

I also made changes to the data by feature engineering to increase signal. This greatly increased the performance of the model, but is not an architectural or parametric change since it is not a change directly to the neural network model itself.

Note: while doing experimenting I did not set random seed. I went back and set a random seed and a few of the results changed. Unfortunately I did not have time to redo all my results with a random seed set.

## 1.2 Comparison of the result with the original output

Model 1 results:

```
Macro F1: 0.5871
Weighted F1: 0.5958
Additional metrics in Appendix 2
```

Final model results:

```
Macro F1: 0.7270
Weighted F1: 0.7424
Additional metrics in Appendix 2
```

Looking at the results we can see that the macro F1 and weighted F1 scores are greatly increased in the final model. Additionally, the final model also greatly improved the precision of all three classes, particularly the neutral class. However, the final model still struggled to discriminate between the neutral class and the two other classes and could use further tuning and feature engineering to make these features more linearly separable. The top 5 stocks (see appendix 2) are also different, however SNDK appears in the output of both stocks suggesting that the starting model still can pick up some signal from stocks that are very likely to outperform – at the time of writing SNDK is up 11.89% in the trading day after training this data. The average percentage gain for the top 5 stocks picked is 6.96%, while the benchmark S&P500 has a gain of 1.54%. However, this is obviously a small sample size and likely is simply because the neural network heavily weighs stocks that have high current momentum and are therefore most likely to increase in value in the near-term.

## 1.3 URL of your Colab code

Code is uploaded into github instead as colab was not used. The code can be run sequentially, although for the comparison the initial data cleaning and test-train split must be run again in order to get the performance of the first model as it was initially trained with less features and fewer performance metrics.

[https://github.com/btorok-bu/METCS767\\_hw2/blob/main/cs767\\_btorok\\_hw2.ipynb](https://github.com/btorok-bu/METCS767_hw2/blob/main/cs767_btorok_hw2.ipynb)

## **2. New Neural Net Application**

### **2.1 Description of the application**

This neural network uses a sequential neural network model and aims to predict stock performance (outperform, neutral, underperform) relative to an S&P500 benchmark performance. The data is sourced from yahoo finance and the massive API. The data contains current stock fundamental statistics, and snapshot market activity from 4 years ago, 2 years ago, 1 year ago, 6 months ago, 3 months ago, and 1 month ago. Through feature engineering the performance of each stock relative to the baseline in these time ranges is computed, standardized, weighted, then a momentum score is calculated. The stocks are then labeled based on the percentile rank of their momentum score.

The feature engineering is improved by adding surge ratios which indicate high volume. To improve momentum metrics further, features for the distance to the 52-week high in addition to distance from the 50- and 200-day moving averages were added. Additionally, a metric for the ratio of the float size to the shares outstanding was added since stocks that have a high number of shares outstanding are currently being shorted, indicating that current outlook for the stock may be poor.

After training the model the performance is measured using loss, overall accuracy, macro F1 score, weighted F1 score, classification report containing the precision and recall for all three classes, and confusion matrix. The probability of the overperform class is then used to get the margin over other classes. The mean of the margin and probability to overperform is then taken to get a confidence score. Stocks are then sorted by their confidence score and returned alongside key metrics. With further tuning and additional data this neural network can aid in discovering stocks that have the highest probability to outperform a chosen benchmark index in a given term.

### **2.2 Summary of your architecture tuning**

For my architecture tuning I first built the initial model given in the class example, I then experimented by changing the output layer to use softmax activation and output 3 results for the 3 classes. I then experimented with adding two additional hidden layers. I also experimented with different activation functions and adding a kernel initializer and then adding an L2 kernel regularizer. I also experimented with trying different optimizers like Adam, Nadam, and SGD optimizers. I found the performance was best with the Adam optimizer with a manual learning rate set, ReLU activation function, 4 layers, with L2 regularizer.

### **2.3 Summary of your hyperparameter tuning**

For my hyperparameter tuning I tried manual learning rates and found that higher learning rates resulted in better performance for my model with the best learning rate being 0.0005. I also found that the best performance for each layer was 512, 256, 128, and 128. The best dropout value was 0.10, 0.10, 0.10 and 0.05 for the final layer. The best L2 regularizer parameter was 8e-4.

## 2.4 Three illustrative input/output pairs from running the implementation

1) creating distance from 50 moving average, distance from 200 moving average, and moving average cross features. This feature provides another metric of the direction of the stock movement and whether the moving averages have been crossed and in what direction.

Input:

```
currentPrice  fiftyDayAverage  twoHundredDayAverage
```

166.50	157.2214	150.05576
65.99	71.1404	68.28100
123.62	131.4586	130.89246
218.04	222.3706	199.54900
250.10	246.7902	296.32050

Output:

```
currentPrice  fiftyDayAverage  twoHundredDayAverage  dist_ma50  dist_ma200  ma_cross
```

166.50	157.2214	150.05576	0.059016	0.109588	0.047753
65.99	71.1404	68.28100	-0.072398	-0.033553	0.041877
123.62	131.4586	130.89246	-0.059628	-0.055561	0.004325
218.04	222.3706	199.54900	-0.019475	0.092664	0.114366
250.10	246.7902	296.32050	0.013411	-0.155981	-0.167151

2) Use the model results to get the probability of being in the overperform class, then get the margin that the p\_over is greater than the greatest of the next class. The mean of these two is used to get the confidence. This way high confidence is only given for stocks with high p\_over and high margin over next.

Input:

```
symbol  p_over  margin_over_vs_next
```

604	CYTK	0.998080	0.996282
345	NVDA	0.966754	0.960818
1397	SNDK	0.915854	0.904881
398	HODD	0.905623	0.890312
643	FIX	0.903497	0.886419

Output:

```
symbol  confidence
```

604	CYTK	0.997181
345	NVDA	0.963786
1397	SNDK	0.910367
398	HODD	0.897967
643	FIX	0.894958

3) Input is the performance metrics from each epoch, the output is the learning curve, using just accuracy, loss, and val\_accuracy

Input: metrics from epoch in appendix 3

Output: output is figure in appendix 3

## 2.5 Key code snippets, with explanation

```
model_25 = tf.keras.models.Sequential([
```

```

# added l2 regularizer with manual tuning of reg param
tf.keras.layers.Flatten(input_shape=(X_train_proc.shape[1],)),
tf.keras.layers.Dense(512, activation='relu',
                      kernel_regularizer=tf.keras.regularizers.l2(8e-4)),
tf.keras.layers.Dropout(0.10),
tf.keras.layers.Dense(256, activation='relu',
                      kernel_regularizer=tf.keras.regularizers.l2(8e-4)),
tf.keras.layers.Dropout(0.10),
tf.keras.layers.Dense(128, activation='relu',
                      kernel_regularizer=tf.keras.regularizers.l2(8e-4)),
tf.keras.layers.Dropout(0.10),
tf.keras.layers.Dense(128, activation='relu',
                      kernel_regularizer=tf.keras.regularizers.l2(8e-4)),
tf.keras.layers.Dropout(0.05)
])

# put together NN with training process, loss, and means of evaluation
# use loss of sparse categorical crossentropy since Dense is in layer
# adjust learning rate in optimizer manually to have slower learning rate
optimizer = tf.keras.optimizers.Adam(learning_rate=0.0001)
loss_fn = tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True)
model_25.compile(optimizer=optimizer,
                  loss=loss_fn,
                  metrics=[metric.SparseCategoricalAccuracy(name='acc'),
                           metric.SparseTopKCategoricalAccuracy(k=2, name='top2_acc'),
                           # per class precision and recall
                           metric.Precision(name='prec_under', class_id=0),
                           metric.Recall(name='rec_under', class_id=0),
                           metric.Precision(name='prec_neutral', class_id=1),
                           metric.Recall(name='rec_neutral', class_id=1),
                           metric.Precision(name='prec_over', class_id=2),
                           metric.Recall(name='rec_over', class_id=2)
                           ])
predictions = model_25.predict(X_train_proc[:1])
tf.nn.softmax(predictions)

model_25.fit(X_train_proc, y_train, epochs=20, verbose=0)

```

The major adjustments here are to change the dropout parameters, add additional layers, add regularization, add additional metrics for each class, and adjust the learning rate for the Adam optimizer. The epochs were also increased but this did not make much of a difference.

Note: A strong starting AI prompt included in appendix 4 was used to consistently excellent AI responses that ensured the code produced did not use different neural network models to get better performance with this dataset.

## 2.6 URL of your code

[https://github.com/btorok-bu/METCS767\\_hw2/tree/main](https://github.com/btorok-bu/METCS767_hw2/tree/main)

## Evaluation

**Note:** We focus the **quality** (not quantity) of your submission. Material that doesn't respond to the criteria reduces grades.

<b>Criterion</b> (based on the value of your significant prompts and your text)	<b>D</b>	<b>C</b>	<b>B</b>	<b>A</b>	<b>Letter Grade</b>	<b>%</b>
<i>Extent to which you added technical content is functional and accurate, and shows demonstrably that you understand the AI output, and have learned technically.</i>	<i>Low extent.</i>	<i>Satisfactory extent</i>	<i>Good extent</i>	<i>Went significantly beyond what's required.</i>		0.0
<i>Extent to which every significant claim you made is supported with clear, relevant logical reasoning that explains its validity.</i>	<i>Low extent.</i>	<i>Satisfactory extent</i>	<i>Good extent</i>	<i>Went significantly beyond what's required.</i>		0.0
<i>Extent to which your added material probes the key mechanisms in depth.</i>	<i>Low extent.</i>	<i>Satisfactory extent</i>	<i>Good extent</i>	<i>Went significantly beyond what's required.</i>		0.0
						Assignment Grade: 0.0

*The resulting grade is the average of these, using A+=100 (outstanding–rare), A=95 (excellent in all ways), A-=90 (excellent), B+=87 (excellent / very good), B=85 (very good), B-=80 (good) etc.*

## Appendix 1

Link to AI chat for getting data: <https://chatgpt.com/share/690fb2bf-acf8-8006-ab93-267e4f24c916>

Link to AI chat for neural network training: <https://chatgpt.com/share/6912ca0e-d138-8006-9edb-6df95f567094>

Note: a separate AI rules document was used for both chats. The AI rules document was referred throughout the conversation by the LLM.

## Appendix 2

First model additional metrics:

	precision	recall	f1-score	support
under	0.67	0.62	0.64	104
neutral	0.39	0.63	0.48	83
over	0.81	0.52	0.63	113
accuracy			0.58	300
macro avg	0.63	0.59	0.59	300
weighted avg	0.65	0.58	0.60	300

Confusion matrix:

[[64 36 4]

[21 52 10]  
[10 44 59]]

Top 5 stocks:

	symbol	p_over	confidence	recommendationKey	currentPrice	forwardPE
604	CYTK	0.982289	0.977802	buy	63.59	-11.797774
345	NVDA	0.935958	0.927752	strong_buy	202.49	49.148060
643	FIX	0.903085	0.892260	strong_buy	965.58	56.832256
1397	SNDK	0.880954	0.869852	buy	199.33	19.709223
51	AZ0	0.881337	0.846761	buy	3674.43	21.078648

Final model additional metrics:

	precision	recall	f1-score	support
under	0.79	0.76	0.77	104
neutral	0.54	0.58	0.56	83
over	0.86	0.84	0.85	113
accuracy			0.74	300
macro avg	0.73	0.73	0.73	300
weighted avg	0.75	0.74	0.74	300

Confusion matrix:

```
[[79 24  1]
 [20 48 15]
 [ 1 17 95]]
```

Top 5 stocks:

398	HOOD	1.000000	1.000000	buy	146.78	201.068480
604	CYTK	1.000000	1.000000	buy	63.59	-11.797774
1397	SNDK	1.000000	1.000000	buy	199.33	19.709223
358	PLTR	1.000000	1.000000	hold	200.47	426.531920
314	MU	1.000000	1.000000	buy	223.77	17.386948

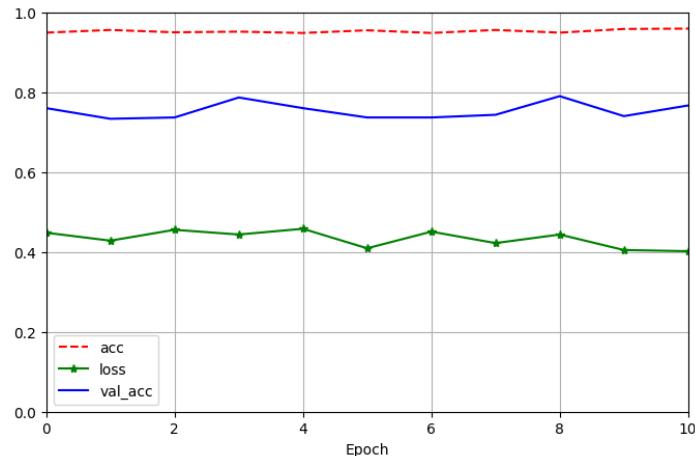
## Appendix 3

Input epoch metrics (only one epoch shown):

Epoch 1/20

```
38/38 1s 13ms/step - acc: 0.9492 - loss: 0.4347 - prec_neutral: 0.7151 - prec_over: 0.6916 - prec_under: 0.5267 - rec_neutral: 0.7088 - rec_over: 0.5465 - rec_under: 0.4736 - top2_acc: 0.9841 - val_acc: 0.7767 - val_loss: 1.1724 - val_prec_neutral: 0.5417 - val_prec_over: 0.6524 - val_prec_under: 0.8453 - val_rec_neutral: 0.7312 - val_rec_over: 0.5245 - val_rec_under: 0.5977 - val_top2_acc: 0.9800
```

Output graph:



The above graph shows that the performance of the neural network does not improve after the first 1-2 epochs. The model converged very fast so improvements can only be made through more feature engineering or further tuning and adjusting neural network architecture.

## Appendix 4

Initial prompt:

Given the ai\_rules.txt document included and this code I have currently, how should I build my label column? I would like to label stocks as outperform, neutral, or underperform using data from the given time ranges for each stock. I unfortunately do not have all fundamental data for each time range, and only have the price history and volume for the time ranges and fundamental data for the current time.

### **Ai\_rules.txt uploaded at the start of the conversation:**

This project is to use a sequential neural network on a dataset containing S&P1500 stocks and snapshot market volume, and close price from 4 years ago, 2 years ago, 1 year ago, 6 months ago, 3 months ago, and 1 month ago. I have the fundamental stock data for the current date.

Instructions:

- act as a senior python engineer
- use concise, production-quality python code
- when suggesting changes to code first attempt to use existing code before generating new code
- you are allowed to suggest changes to existing code in order to make new functions work better
- always explain the reasoning before showing code
- when code may be complex, add inline comments to explain code
- do more explaining in the prompt response than in code comments
- never use any emojis in responses including code comments
- prioritize pandas when making changes to the existing data
- follow PEP8 standards for generated code
- do not suggest any module for neural networks outside of tensorflow and do not suggest more complex models
- always refer back to this document before answering.

These are all of the column headers in the CSV I will be using for analysis.

address1	city	state	zip	country	phone	website	industry
industryKey	industryDisp	sector	sectorKey	sectorDisp			
longBusinessSummary		fullTimeEmployees	companyOfficers		auditRisk		
boardRisk	compensationRisk	shareHolderRightsRisk		overallRisk			
governanceEpochDate		compensationAsOfEpochDate		irWebsite			
executiveTeam	maxAge	priceHint	previousClose	open	dayLow		
dayHigh	regularMarketPreviousClose	regularMarketOpen					
regularMarketDayLow	regularMarketDayHigh		dividendRate	dividendYield			
exDividendDate	payoutRatio	fiveYearAvgDividendYield	beta	trailingPE			
forwardPE	volumeregularMarketVolume	averageVolume					
averageVolume10days	averageDailyVolume10Day	bid	ask	bidSize			
askSize	marketCap	fiftyTwoWeekLow	fiftyTwoWeekHigh	allTimeHigh			
allTimeLow	priceToSalesTrailing12Months		fiftyDayAverage				
twoHundredDayAverage	trailingAnnualDividendRate	trailingAnnualDividendYield					
currency	tradeable	enterpriseValue	profitMargins	floatShares			
sharesOutstanding	sharesShort	sharesShortPriorMonth					
sharesShortPreviousMonthDate	dateShortInterest	sharesPercentSharesOut					
heldPercentInsiders	heldPercentInstitutions		shortRatio	shortPercentOfFloat			
impliedSharesOutstanding	bookValue	priceToBook	lastFiscalYearEnd				
nextFiscalYearEnd	mostRecentQuarter	earningsQuarterlyGrowth					
netIncomeToCommon	trailingEps	forwardEps	lastSplitFactor	lastSplitDate			
enterpriseToRevenue	enterpriseToEbitda	52WeekChange	SandP52WeekChange				
lastDividendValue	lastDividendDate	quoteType	currentPrice				
targetHighPrice	targetLowPrice	targetMeanPrice	targetMedianPrice				
recommendationMean	recommendationKey	numberOfAnalystOpinions	totalCash				
totalCashPerShare	ebitda	totalDebt	quickRatio	currentRatio	totalRevenue		
debtToEquity	revenuePerShare	returnOnAssets	returnOnEquity				
grossProfits	freeCashflow	operatingCashflow	earningsGrowth				
revenueGrowth	grossMargins	ebitdaMargins	operatingMargins				
financialCurrency	symbol	language	region	typeDisp	quoteSourceName		
triggerable	customPriceAlertConfidence	marketState		earningsTimestamp			
earningsTimestampStart	earningsTimestampEnd		earningsCallTimestampStart				
earningsCallTimestampEnd	isEarningsDateEstimate		epsTrailingTwelveMonths				
epsForward	epsCurrentYear	priceEpsCurrentYear	fiftyDayAverageChange				
fiftyDayAverageChangePercent	twoHundredDayAverageChange						
twoHundredDayAverageChangePercent	sourceInterval	exchangeDataDelayedBy					
averageAnalystRating	cryptoTradeable	shortName	longName				
hasPrePostMarketData	firstTradeDateMilliseconds	postMarketChangePercent					

postMarketPrice	postMarketChange	regularMarketChange				
regularMarketDayRange	fullExchangeName	averageDailyVolume3Month				
fiftyTwoWeekLowChange	fiftyTwoWeekLowChangePercent	fiftyTwoWeekRange				
fiftyTwoWeekHighChange	fiftyTwoWeekHighChangePercent					
fiftyTwoWeekChangePercent	dividendDate	regularMarketChangePercent				
corporateActions	postMarketTime	regularMarketTime	exchange			
messageBoardId	exchangeTimezoneName	exchangeTimezoneShortName				
gmtOffsetMilliseconds	market	esgPopulated	regularMarketPrice			
trailingPegRatio	address2	displayName	fax	ipoExpectedDate		
prevName	nameChangeDate	industrySymbol	prevTicker			
tickerChangeDate	4y_date	4y_open	4y_high	4y_low		
4y_close	4y_volume	4y_vwap	2y_date	2y_open	2y_high	
2y_low	2y_close	2y_volume	2y_vwap	1y_date	1y_open	
1y_high	1y_low	1y_close	1y_volume	1y_vwap	6m_date	
6m_open	6m_high	6m_low	6m_close	6m_volume	6m_vwap	
3m_date	3m_open	3m_high	3m_low	3m_close	3m_volume	
3m_vwap	1m_date	1m_open	1m_high	1m_low	1m_close	
1m_volume	1m_vwap	massive_error				